

# DataFlowTasks.jl

Julia Tasks which automatically handle data-dependencies

Luiz M. Faria<sup>1</sup>    François Févotte<sup>2</sup>

<sup>1</sup>Research scientist  
INRIA

<sup>2</sup>Chief Scientist  
TriScale innov


JuliaCon Local Eindhoven  
December 1, 2023




# Overview

- DataFlowTasks.jl is a Julia package dedicated to parallel programming on **multi-core shared memory CPUs**.
- Automatically infer Task interdependencies based on user annotations (`@R`, `@W`, `@RW`).
- Inspired by task programming libraries such as StarPU / PaRSEC
  - (in spirit, not in ambition)
- Simple API: `@dspawn` macro.

```
function foo!(A, b)
    fill!(b, 0)
    view(b, 1:2) .+= 2
    view(b, 3:4) .+= 3
    A \ b
end
```



```
function foo!(A, b)
    @dspawn fill!(@W(b), 0)           # task 1
    @dspawn @RW(view(b, 1:2)) .+= 2  # task 2
    @dspawn @RW(view(b, 3:4)) .+= 3  # task 3
    @dspawn @R(A) \ @R(b)             # task 4
end
```



# Table of Contents

- 1 Introduction and high-level usage
  - Task based parallelism in Julia
  - Motivation
  - Basic idea
- 2 Real use case: tiled Cholesky
  - Algorithm
  - Sequential & parallel implementation
  - Profiling & debugging tools
  - Comparison to OpenBLAS
- 3 Concluding remarks

# Table of Contents


- 1 Introduction and high-level usage
  - Task based parallelism in Julia
  - Motivation
  - Basic idea
- 2 Real use case: tiled Cholesky
  - Algorithm
  - Sequential & parallel implementation
  - Profiling & debugging tools
  - Comparison to OpenBLAS
- 3 Concluding remarks

# Task based parallelism

- A task is a unit of execution or a unit of work
- **Task** objects can be created using `@task`
- Once created, **Task** objects must be scheduled for execution
- Usually, **Tasks** are created + scheduled using `@spawn`
- Responsibility of synchronizing tasks is left to the programmer

## Example (Sequential)

```
# Task
b = zeros(4)           # 1 - Initialization
b[1:2] .+= 2           # 2 - Work on first half
b[3:4] .+= 3           # 3 - Work on second half
res = A \ b            # 4 - Use entire vector
#
```




# Task based parallelism

- A task is a unit of execution or a unit of work
- **Task** objects can be created using `@task`
- Once created, **Task** objects must be scheduled for execution
- Usually, **Tasks** are created + scheduled using `@spawn`
- **Responsibility of synchronizing tasks is left to the programmer**

## Example (Synchronizing tasks with return values)

```
                                # Task
t1 = @spawn zeros(4)           # 1
t2 = @spawn (b = fetch(t1); b[1:2] .+ 2) # 2
t3 = @spawn (b = fetch(t1); b[3:4] .+ 3) # 3
t4 = @spawn A \ vcat(fetch(t2), fetch(t3)) # 4
fetch(t4) # get result
```




# Task based parallelism

- A task is a unit of execution or a unit of work
- **Task** objects can be created using **@task**
- Once created, **Task** objects must be scheduled for execution
- Usually, **Tasks** are created + scheduled using **@spawn**
- **Responsibility of synchronizing tasks is left to the programmer**

## Example (Synchronizing tasks with explicit barriers)

```
b = Vector{Float64}(undef, 4)           # Task
t1 = @spawn fill!(b, 0)                  # 1
t2 = @spawn (wait(t1); view(b,1:2) += 2) # 2
t3 = @spawn (wait(t1); view(b,3:4) += 3) # 3
t4 = @spawn (wait(t2); wait(t3); A \ b)  # 4
fetch(t4) # get result
```



# Motivation for DataFlowTasks

- Reasoning about **Task** interdependencies can be challenging
- Specially for algorithms making constant re-use of data
- Sometimes, it is simpler to reason about how **Tasks** depend on data than how **Tasks** depend on each other

## Example (Synchronizing tasks)

```
b = rand(4)
t1 = @spawn fill!(b,0)           # RW access to b
t2 = @spawn (wait(t1); view(b,1:2) += 2) # RW access to b[1:2]
t3 = @spawn (wait(t1); view(b,3:4) += 3) # RW access to b[3:4]
t4 = @spawn (wait(t2); wait(t3); A \ b) # R access to A and b
fetch(t4)
```



# Motivation for DataFlowTasks

- Reasoning about **Task** interdependencies can be challenging
- Specially for algorithms making constant re-use of data
- Sometimes, it is simpler to reason about how **Tasks** depend on data than how **Tasks** depend on each other

## Example (Synchronizing tasks)

```
b = rand(4)
t1 = @spawn fill!(b,0)           # RW access to b
t2 = @spawn (wait(t1); view(b,1:2) += 2) # RW access to b[1:2]
t3 = @spawn (wait(t1); view(b,3:4) += 3) # RW access to b[3:4]
t4 = @spawn (wait(t2); wait(t3); A \ b) # R access to A and b
fetch(t4)
```

- Would like to declare only the task-to-data dependencies

# Basic idea

- 1 Extract data dependency from user annotations
- 2 Infer task dependency from data dependency
- 3 Schedule tasks with inferred dependencies


b

## User-written code

```
b = rand(4)
```

# Basic idea

- 1 Extract data dependency from user annotations
- 2 Infer task dependency from data dependency
- 3 Schedule tasks with inferred dependencies




## User-written code

```
b = rand(4)
@dspawn fill!(@W(b),0)
```

# Basic idea

- 1 Extract data dependency from user annotations
- 2 Infer task dependency from data dependency
- 3 Schedule tasks with inferred dependencies




## User-written code

```
b = rand(4)
@spawn fill!(@W(b),0)
@spawn @RW(view(b,1:2)) += 2
```

# Basic idea

- 1 Extract data dependency from user annotations
- 2 Infer task dependency from data dependency
- 3 Schedule tasks with inferred dependencies



## User-written code


```
b = rand(4)
@spawn fill!(@W(b), 0)
@spawn @RW(view(b, 1:2)) += 2
```

## Behind the scenes: DAG

```
# Note the quadratic complexity!
for i in 1:N, j in i:-1:1
    # Detect conflict between i and j
    for di in data(i), dj in data(j)
        if memory_overlap(di, dj)
            add_edge(j, i)
```

# Basic idea

- 1 Extract data dependency from user annotations
- 2 Infer task dependency from data dependency
- 3 Schedule tasks with inferred dependencies



## User-written code


```
b = rand(4)
@spawn fill!(@W(b), 0)
@spawn @RW(view(b, 1:2)) += 2
@spawn @RW(view(b, 3:4)) += 3
```

## Behind the scenes: DAG

```
# Note the quadratic complexity!
for i in 1:N, j in i:-1:1
    # Detect conflict between i and j
    for di in data(i), dj in data(j)
        if memory_overlap(di, dj)
            add_edge(j, i)
```

# Basic idea

- 1 Extract data dependency from user annotations
- 2 Infer task dependency from data dependency
- 3 Schedule tasks with inferred dependencies



## User-written code


```
b = rand(4)
@spawn fill!(@W(b), 0)
@spawn @RW(view(b, 1:2)) += 2
@spawn @RW(view(b, 3:4)) += 3
```

## Behind the scenes: DAG

```
# Note the quadratic complexity!
for i in 1:N, j in i:-1:1
    # Detect conflict between i and j
    for di in data(i), dj in data(j)
        if memory_overlap(di, dj)
            add_edge(j, i)
```

## Basic idea

- 1 Extract data dependency from user annotations
- 2 Infer task dependency from data dependency
- 3 Schedule tasks with inferred dependencies



## User-written code

```
b = rand(4)
@spawn fill!(@W(b),0)
@spawn @RW(view(b,1:2)) += 2
@spawn @RW(view(b,3:4)) += 3
t4 = @spawn A \ @R(b)
fetch(t4)
```


## Behind the scenes: DAG

```
# Note the quadratic complexity!
for i in 1:N, j in i:-1:1
    # Detect conflict between i and j
    for di in data(i), dj in data(j)
        if memory_overlap(di, dj)
            add_edge(j, i)
```



# Basic idea

- 1 Extract data dependency from user annotations
- 2 Infer task dependency from data dependency
- 3 Schedule tasks with inferred dependencies



## User-written code


```
b = rand(4)
@spawn fill!(@W(b),0)
@spawn @RW(view(b,1:2)) += 2
@spawn @RW(view(b,3:4)) += 3
t4 = @spawn A \ @R(b)
fetch(t4)
```

## Behind the scenes: DAG

```
# Note the quadratic complexity!
for i in 1:N, j in i:-1:1
    # Detect conflict between i and j
    for di in data(i), dj in data(j)
        if memory_overlap(di, dj)
            add_edge(j, i)
```

# Basic idea

- 1 Extract data dependency from user annotations
- 2 Infer task dependency from data dependency
- 3 Schedule tasks with inferred dependencies



## User-written code


```
b = rand(4)
@spawn fill!(@W(b), 0)
@spawn @RW(view(b, 1:2)) += 2
@spawn @RW(view(b, 3:4)) += 3
t4 = @spawn A \ @R(b)
fetch(t4)
```

## Behind the scenes: DAG

```
# Note the quadratic complexity!
for i in 1:N, j in i:-1:1
    # Detect conflict between i and j
    for di in data(i), dj in data(j)
        if memory_overlap(di, dj)
            add_edge(j, i)
        end
    end
end
```

# Basic idea

- 1 Extract data dependency from user annotations
- 2 Infer task dependency from data dependency
- 3 Schedule tasks with inferred dependencies



## User-written code

```
b = rand(4)
@spawn fill!(@W(b),0)
@spawn @RW(view(b,1:2)) += 2
@spawn @RW(view(b,3:4)) += 3
t4 = @spawn A \ @R(b)
fetch(t4)
```

## Behind the scenes: task scheduling

```
t4 = Threads.@spawn begin
    wait(t2); wait(t3)
    A \ b
end
```

# Table of Contents

- 1 Introduction and high-level usage
  - Task based parallelism in Julia
  - Motivation
  - Basic idea
- 2 Real use case: tiled Cholesky
  - Algorithm
  - Sequential & parallel implementation
  - Profiling & debugging tools
  - Comparison to OpenBLAS
- 3 Concluding remarks

# Tiled Cholesky factorization: idea

- Objective: take a SPD matrix  $A$ , find lower triangular  $L$  such that

$$A = LL^T$$

- Idea for a tiled algorithm:

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{12}^T & A_{22} \end{bmatrix} = \begin{bmatrix} L_{11} & \\ L_{21} & L_{22} \end{bmatrix} \begin{bmatrix} L_{11}^T & L_{21}^T \\ & L_{22}^T \end{bmatrix}$$

# Tiled Cholesky factorization: idea

- Objective: take a SPD matrix  $A$ , find lower triangular  $L$  such that

$$A = LL^\top$$

- Idea for a tiled algorithm:

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{12}^\top & A_{22} \end{bmatrix} = \begin{bmatrix} L_{11} & \\ L_{21} & L_{22} \end{bmatrix} \begin{bmatrix} L_{11}^\top & L_{21}^\top \\ & L_{22}^\top \end{bmatrix}$$

①  $A_{11} = L_{11}L_{11}^\top$

▷ Cholesky factorization to get  $L_{11}$

cholesky!

# Tiled Cholesky factorization: idea

- Objective: take a SPD matrix  $A$ , find lower triangular  $L$  such that

$$A = LL^T$$

- Idea for a tiled algorithm:

$$\begin{bmatrix} A_{11} & A_{12} \\ A_{12}^T & A_{22} \end{bmatrix} = \begin{bmatrix} L_{11} & \\ L_{21} & L_{22} \end{bmatrix} \begin{bmatrix} L_{11}^T & L_{21}^T \\ & L_{22}^T \end{bmatrix}$$

①  $A_{11} = L_{11}L_{11}^T$

▷ Cholesky factorization to get  $L_{11}$

cholesky!

②  $A_{12} = L_{11}L_{21}^T$

▷ Triangular solve to get  $L_{21}$

ldiv!

# Tiled Cholesky factorization: idea

- Objective: take a SPD matrix  $A$ , find lower triangular  $L$  such that

$$A = LL^T$$

- Idea for a tiled algorithm:


$$\begin{bmatrix} A_{11} & A_{12} \\ A_{12}^T & A_{22} \end{bmatrix} = \begin{bmatrix} L_{11} & \\ L_{21} & L_{22} \end{bmatrix} \begin{bmatrix} L_{11}^T & L_{21}^T \\ & L_{22}^T \end{bmatrix}$$

- 1  $A_{11} = L_{11}L_{11}^T$   
▷ Cholesky factorization to get  $L_{11}$  cholesky!
- 2  $A_{12} = L_{11}L_{21}^T$   
▷ Triangular solve to get  $L_{21}$  ldiv!
- 3  $A_{22} - L_{21}L_{21}^T = L_{22}L_{22}^T$   
▷ Matmul to get  $A_{22} - L_{21}L_{21}^T$  mul!  
▷ Cholesky factorization to get  $L_{22}$  cholesky!




# Tiled Cholesky factorization: algorithm

- Input: symmetric positive-definite matrix  $A$  with  $N \times N$  blocks
- Goal: compute lower-triangular  $L$  such that  $A = LL^\top$



# Tiled Cholesky factorization: algorithm

- Input: symmetric positive-definite matrix  $A$  with  $N \times N$  blocks
- Goal: compute lower-triangular  $L$  such that  $A = LL^\top$




i=1

```
for i in 1:N
  ① cholesky!
     $A_{i,i} \leftarrow \text{chol!}(A_{i,i})$ 
     $L_{i,i} = \text{LowerTriangular}(A_{i,i})$ 
```

# Tiled Cholesky factorization: algorithm

- Input: symmetric positive-definite matrix  $A$  with  $N \times N$  blocks
- Goal: compute lower-triangular  $L$  such that  $A = LL^\top$




i=1

```
for i in 1:N
  ① cholesky!
     $A_{i,i} \leftarrow \text{chol!}(A_{i,i})$ 
     $L_{i,i} = \text{LowerTriangular}(A_{i,i})$ 
  ② ldiv!
    for k in 2:N
       $A_{k,i} \leftarrow \text{ldiv!}(L_{i,i}, A_{k,i}),$ 
```

# Tiled Cholesky factorization: algorithm

- Input: symmetric positive-definite matrix  $A$  with  $N \times N$  blocks
- Goal: compute lower-triangular  $L$  such that  $A = LL^\top$




$i=1$

```
for i in 1:N
  ① cholesky!
     $A_{i,i} \leftarrow \text{chol!}(A_{i,i})$ 
     $L_{i,i} = \text{LowerTriangular}(A_{i,i})$ 
  ② ldiv!
    for k in 2:N
       $A_{k,i} \leftarrow \text{ldiv!}(L_{i,i}, A_{k,i}),$ 
  ③ mul!
    for k in i+1:N, j in i+1:k
       $A_{k,j} \leftarrow A_{k,j} - A_{k,i}A_{i,j}^\top$ 
```

# Tiled Cholesky factorization: algorithm

- Input: symmetric positive-definite matrix  $A$  with  $N \times N$  blocks
- Goal: compute lower-triangular  $L$  such that  $A = LL^\top$




$i=2$ , repeat

```
for i in 1:N
  ① cholesky!
     $A_{i,i} \leftarrow \text{chol!}(A_{i,i})$ 
     $L_{i,i} = \text{LowerTriangular}(A_{i,i})$ 
  ② ldiv!
    for k in 2:N
       $A_{k,i} \leftarrow \text{ldiv!}(L_{i,i}, A_{k,i}),$ 
  ③ mul!
    for k in i+1:N, j in i+1:k
       $A_{k,j} \leftarrow A_{k,j} - A_{k,i}A_{k,i}^\top$ 
```

# Tiled Cholesky factorization: algorithm


- Input: symmetric positive-definite matrix  $A$  with  $N \times N$  blocks
- Goal: compute lower-triangular  $L$  such that  $A = LL^\top$



```
for i in 1:N
  ① cholesky!
     $A_{i,i} \leftarrow \text{chol!}(A_{i,i})$ 
     $L_{i,i} = \text{LowerTriangular}(A_{i,i})$ 
  ② ldiv!
    for k in 2:N
       $A_{k,i} \leftarrow \text{ldiv!}(L_{i,i}, A_{k,i}),$ 
  ③ mul!
    for k in i+1:N, j in i+1:k
       $A_{k,j} \leftarrow A_{k,j} - A_{k,i}A_{k,i}^\top$ 
```

- Output:  $L$  in the lower-triangular part of  $A$

# Tiled Cholesky factorization: complexity



For an  $N \times N$  block matrix:

- $\mathcal{O}(N)$  **cholesky factorizations**
- $\mathcal{O}(N^2)$  **triangular solves**
- $\mathcal{O}(N^3)$  **matrix multiplications**

Globally compute-bound:

- $\mathcal{O}(N^3)$  flops for
  - $\mathcal{O}(N^2)$  bytes
- ⇒ potential for efficient parallelism

# Tiled Cholesky factorization: sequential implementation

```
function cholesky_tiled!(A, ts)
    m = size(A, 1); n = m ÷ ts      # Matrix size & number of tiles
    T = [view(A, tilerange(i, ts), tilerange(j, ts)) for i in 1:n, j in 1:n]

    for i in 1:n
        cholesky!(T[i,i])           # Diagonal cholesky serial factorization

        U = UpperTriangular(T[i,i]) # Left tiles update
        for j in i+1:n
            ldiv!(U', T[i,j])
        end

        for j in i+1:n, k in j:n     # Submatrix update
            mul!(T[j,k], T[i,j]', T[i,k], -1, 1)
        end
    end

    # Construct the factorized object
    return Cholesky(A, 'U', zero(LinearAlgebra.BlasInt))
end
```



# Tiled Cholesky factorization: parallel implementation

```
function cholesky_dft!(A, ts)
    m = size(A, 1); n = m ÷ ts           # Matrix size & number of tiles
    T = [view(A, tilerange(i, ts), tilerange(j, ts)) for i in 1:n, j in 1:n]

    for i in 1:n
        @dspawn cholesky!(@RW(T[i,i]))    label="chol ($i,$i)"

        U = UpperTriangular(T[i,i])
        for j in i+1:n
            @dspawn ldiv!(@R(U)', @RW(T[i,j]))    label="ldiv ($i,$j)"
        end

        for j in i+1:n, k in j:n
            @dspawn mul!(@RW(T[j,k]), @R(T[i,j])', @R(T[i,k]), -1, 1) label="schur($j,$k)"
        end
    end

    # Construct the factorized object
    r = @dspawn Cholesky(@R(A), 'U', zero(LinearAlgebra.BlasInt)) label="result"
    return fetch(r)
end
```

# Tiled Cholesky factorization: performance

```
julia> using BenchmarkTools
```

```
julia> n = 4096; ts = 512;
```

```
julia> t_seq = @belapsed(cholesky_tiled!(A, ts),  
                        setup=(A=spd_matrix(n)), evals=1)  
0.559141485
```

```
julia> t_par = @belapsed(cholesky_dft!(A, ts),  
                        setup=(A=spd_matrix(n)), evals=1)  
0.108479304
```

```
julia> (; threads = Threads.nthreads(),  
        speedup = t_seq / t_par)  
(threads = 8, speedup = 5.15436089063056)
```

# Tiled Cholesky factorization: profiling & debugging

```
julia> n = 4096; ts = 512;
```

```
julia> A = spd_matrix(n);
```

```
julia> log_info = DataFlowTasks.@log cholesky_dft!(Ac, ts)
```


LogInfo with 121 logged tasks

```
julia> DataFlowTasks.describe(log_info; categories=["chol", "ldiv", "schur"])
```

- Elapsed time : 0.297
  - + Critical Path : 0.253
  - + No-Wait : 0.188
  
- Run time : 2.379
  - + Computing : 1.507
    - o chol : 0.061
    - o ldiv : 0.414
    - o schur : 1.032
    - o unlabeled : 0.000
  - + Task Insertion : 0.003
  - + Other (idle) : 0.870





# Tiled Cholesky factorization: profiling & debugging

```
julia> DataFlowTasks.stack_weakdeps_env!()
julia> using GraphViz
[ Info: Loading DataFlowTasks dag plot utilities
julia> GraphViz.Graph(log_info)
```



# Tiled Cholesky factorization: profiling & debugging



```
julia> using CairoMakie # or GLMakie for more interactivity
[ Info: Loading DataFlowTasks profile plot utilities
julia> plot(log_info; categories=["chol", "ldiv", "schur"])
```




How does it compare to multi-threaded OpenBLAS?

# Tiled Cholesky factorization: performance

Cholesky factorization on 8 threads



# Tiled Cholesky factorization: performance





# Table of Contents

- 1 Introduction and high-level usage
  - Task based parallelism in Julia
  - Motivation
  - Basic idea
- 2 Real use case: tiled Cholesky
  - Algorithm
  - Sequential & parallel implementation
  - Profiling & debugging tools
  - Comparison to OpenBLAS
- 3 Concluding remarks

## Concluding remarks

- DataFlowTasks.jl tries to make parallel programming easier for *certain types of algorithms*
- Simple API: `@dspawn` macro
- Several supplementary examples in the online documentation

<https://github.com/maltezfaria/DataFlowTasks.jl>

If you think DataFlowTasks.jl may help parallelize your algorithm,  
please give it a try or reach out to us!

# Appendix

# Table of Contents

## 4 Implementation details

- DataFlowTask
- DAG
- TaskGraph
- Logging

# Main types

The package revolves around three main types:

- `DataFlowTask`: wrapper around a `Task` with user-declared data dependencies

# Main types

The package revolves around three main types:

- `DataFlowTask`: wrapper around a `Task` with user-declared data dependencies
- `DAG`: graph data structure to represent dependencies between `DataFlowTasks`

# Main types

The package revolves around three main types:

- `DataFlowTask`: wrapper around a `Task` with user-declared data dependencies
- `DAG`: graph data structure to represent dependencies between `DataFlowTasks`
- `TaskGraph`: a `DAG` and some helper functions/workers to manage it

# Main types

The package revolves around three main types:

- `DataFlowTask`: wrapper around a `Task` with user-declared data dependencies
- `DAG`: graph data structure to represent dependencies between `DataFlowTasks`
- `TaskGraph`: a `DAG` and some helper functions/workers to manage it

Next: some technical details about their implementation.



## DataFlowTask: @dspawn macro

DataFlowTasks are created using `@dspawn`. The macro does the following:

- Scan the `Expr` for `@R`, `@W`, `@RW` annotations
- Create a data and mode tuples
- Remove annotations from the `Expr`
- Parse keyword arguments
- Create an anonymous function wrapping the new `Expr`
- Insert a call to DataFlowTask constructor

### User-written code

```
@dspawn(foo!(@W(B), @R(A)),  
        label="foo!")
```

### Macro expansion (approx.)

```
_t = DataFlowTask(  
    () -> foo!(B, A),  
    (B, A),  
    (WRITE, READ);  
    label="foo!")  
spawn(_t)
```

# DataFlowTask structure

Inner-constructor of DataFlowTask handles much of the insertion/removal logic:

---

```
mutable struct DataFlowTask
  data::Tuple
  access_mode::NTuple{<:Any, AccessMode}
  task::Task
  ...
  function DataFlowTask(f, data, mode, taskgraph)
    tj = new(data, mode) # incomplete initialization
    addnode!(taskgraph, tj, true)
    deps = inneighbors(taskgraph, tj) |> copy
    tj.task = @task do
      foreach(wait, deps)
      res = f() # run the underlying function
      put!(taskgraph.finished, tj)
      return res
    end
  end
end
end
```

# DAG

Graph structure used to represent dependencies between DataFlowTasks:

- Dynamic: nodes are added and removed on the fly
  - Buffered: limit the number of active nodes
  - Thread-safe: multiple threads can add/remove nodes
  - Efficient: easy access to both in- and out-neighbors
- 

```
struct DAG{T}
  inoutlist::OrderedDict{...}
  cond_push::Condition
  lock::ReentrantLock
  sz_max::Base.RefValue{Int}
  ...
end
```

# DAG

Graph structure used to represent dependencies between DataFlowTasks:

- Dynamic: nodes are added and removed on the fly
  - **Buffered: limit the number of active nodes**
  - Thread-safe: multiple threads can add/remove nodes
  - Efficient: easy access to both in- and out-neighbors
- 

```
struct DAG{T}
  inoutlist::OrderedDict{...}
  cond_push::Condition
  lock::ReentrantLock
  sz_max::Base.RefValue{Int}
  ...
end
```

- ▷ addnode!(dag,node) calls wait(cond\_push) if full
- ▷ removenode!(dag,node) calls notify(cond\_push)

Graph structure used to represent dependencies between DataFlowTasks:

- Dynamic: nodes are added and removed on the fly
  - Buffered: limit the number of active nodes
  - **Thread-safe: multiple threads can add/remove nodes**
  - Efficient: easy access to both in- and out-neighbors
- 

```
struct DAG{T}
  inoutlist::OrderedDict{...}
  cond_push::Condition
  lock::ReentrantLock
  sz_max::Base.RefValue{Int}
  ...
end
```

- ▷ mutating the DAG requires acquiring/releasing the lock
- ▷ pattern: `@lock dag.lock` code

# Taskgraph

Essentially a `DAG{DataFlowTask}` with

- A channel to store finished tasks
- A dedicated `Task` to remove nodes from the graph

---

```
mutable struct TaskGraph
  dag::DAG{DataFlowTask}
  finished::FinishedChannel
  dag_cleaner::Task
  function TaskGraph(sz)
    dag = DAG{DataFlowTask}(sz)
    finished = FinishedChannel()
    tg = new(dag, finished)
    start_dag_cleaner(tg)
    return tg
  end
end
```

- ▷ Insertion done by `DataFlowTask` constructor
- ▷ Removal done in two steps:
  - 1 The `node.task` moves the node into the finished channel
  - 2 Dedicated task handles finished channel

Some logging capabilities available:

- `@log` macro logs the execution of block
- `describe(loginfo)` shows a summary
- `Graph(loginfo)` displays the DAG
- `plot(loginfo)` plots the execution trace

# Logging

Basic idea:

- Redefine the function `should_log()` to control logging
- Tasks conditionally create a `TaskLog` object
- Information dumped into `LogInfo` object
- Logging should have zero overhead when disabled

---

```
struct TaskLog
  tag::Int
  time_start::UInt64
  time_finish::UInt64
  tid::Int
  inneighbors::Vector{Int64}
  label::String
end
```



Known limitations:

- ▷ If the function yields, logged task time is not representative of execution time





# Tiled Cholesky factorization: performance

Cholesky factorization on 8 threads



# Tiled Cholesky factorization: performance

Cholesky factorization on 8 threads



# Tiled Cholesky factorization: performance

